

Social network Graph Link Prediction - Facebook Challenge

Problem statement:

Given a directed social graph, have to predict missing links to recommend users (Link Prediction in graph)

Data Overview

Taken data from facebook's recruiting challenge on kaggle

<https://www.kaggle.com/c/FacebookRecruiting>

data contains two columns source and destination eac edge in graph

```
- Data columns (total 2 columns):  
- source_node          int64  
- destination_node     int64
```

Mapping the problem into supervised learning problem:

- Generated training samples of good and bad links from given directed graph and for each link got some features like no of followers, is he followed back, page rank, katz score, adar index, some svd fetures of adj matrix, some weight features etc. and trained ml model based on these features to predict link.
- Some reference papers and videos :
 - <https://www.cs.cornell.edu/home/kleinber/link-pred.pdf>
 - <https://www3.nd.edu/~dial/publications/lichtenwalter2010new.p>
 - https://kaggle2.blob.core.windows.net/forum-message-attachments/2594/supervised_link_prediction.pdf
 - <https://www.youtube.com/watch?v=2M77Hgy17cg>

Business objectives and constraints:

- No low-latency requirement.
- Probability of prediction is useful to recommend ighest probability links

Performance metric for supervised learning:

- Both precision and recall is important so F1 score is good choice
- Confusion matrix

```
In [1]: #Importing Libraries  
# please do go through this python notebook:  
import warnings  
warnings.filterwarnings("ignore")  
  
import csv  
import pandas as pd #pandas to create small dataframes  
import datetime #Convert to unix time
```

```

import time #Convert to unix time
# if numpy is not installed already : pip3 install numpy
import numpy as np#Do arithmetic operations on arrays
# matplotlib: used to plot graphs
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns#Plots
from matplotlib import rcParams#Size of plots
from sklearn.cluster import MiniBatchKMeans, KMeans#Clustering
import math
import pickle
import os
# to install xgboost: pip3 install xgboost
import xgboost as xgb

import warnings
import networkx as nx
import pdb
import pickle
from pandas import HDFStore, DataFrame
from pandas import read_hdf
from scipy.sparse.linalg import svds, eigs
import gc
from tqdm import tqdm
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score

```

Assignments:

1. Add another feature called Preferential Attachment with followers and followees data of vertex. you can check about Preferential Attachment in below link <http://be.amazd.com/link-prediction/>
2. Add feature called svd_dot. you can calculate svd_dot as Dot product between source node svd and destination node svd features. you can read about this in below pdf https://storage.googleapis.com/kaggle-forum-message-attachments/2594/supervised_link_prediction.pdf
3. Tune hyperparameters for XG boost with all these features and check the error metric.

1. Preferential Attachment

```

In [38]: # estimating how "rich" each vertices are by calculating the number of friends/neighbourhood ( $|\Gamma(x)|$ )

# reading graph using non_direction
directed_graph=nx.read_edgelist('data/after_eda/train_woheader.csv',delimiter=',',create_using=nx.DiGraph(),nodetype=int)

# prfrential attachment of each user_i
def ps_attachment_single_user(a):
    try:
        preferential_score=len(set(directed_graph.successors(a)))
    except:
        return preferential_score

```

```

except:
    return -1

```

```

In [39]: # reading graph using non_direction as inbuilt Networkx.preferential_attachment does not work on directed graph
non_directed_graph=nx.read_edgelist('data/after_eda/train_woheader.csv',delimiter=',',create_using=nx.Graph(),nodetype=int)

# preferential_attachment for (u_i,u_j) using Networkx library to calculate preferential_attachment score
def ps_attachment(a,b):
    try:
        preferential_score=nx.preferential_attachment(non_directed_graph,[(a,b)])
        for s in preferential_score:
            return s[2]

    except:
        return -1

```

```

In [40]: from pandas import read_hdf
df_final_train = read_hdf('data/fea_sample/storage_sample_stage4.h5', 'train_df',mode='r')
df_final_test = read_hdf('data/fea_sample/storage_sample_stage4.h5', 'test_df',mode='r')

if not os.path.isfile('data/fea_sample/storage_sample_stage5.h5'):
    #=====

    # preferential_attachment of source
    ps_attachment_source_train = df_final_train.apply(lambda row: ps_attachment_single_user(row['source_node']),axis=1)
    df_final_train["preferential_score_source"] = ps_attachment_source_train

    ps_attachment_source_test = df_final_test.apply(lambda row: ps_attachment_single_user(row['source_node']),axis=1)
    df_final_test["preferential_score_source"] = ps_attachment_source_test

    # preferential_attachment of destination
    ps_attachment_destination_train = df_final_train.apply(lambda row: ps_attachment_single_user(row['destination_node']),axis=1)
    df_final_train["preferential_score_destination"] = ps_attachment_destination_train

    ps_attachment_destination_test = df_final_test.apply(lambda row: ps_attachment_single_user(row['destination_node']),axis=1)
    df_final_test["preferential_score_destination"] = ps_attachment_destination_test

    # preferential_attachment_score by multiplying size neighbourhood of source and destination(using networkx library)
    preferential_score_source_destination_train = df_final_train.apply(lambda row:
ps_attachment(row['source_node'],row['destination_node']),axis=

```

```

1)
df_final_train["preferential_score_source_destination"] = p
referential_score_source_destination_train

preferential_score_source_destination_test = df_final_test.
apply(lambda row:

ps_attachment(row['source_node'],row['destination_node']),axis=
1)

df_final_test["preferential_score_source_destination"] = pr
eferential_score_source_destination_test

hdf = HDFStore('data/fea_sample/storage_sample_stage5.h5')
hdf.put('train_df',df_final_train, format='table', data_col
umns=True)
hdf.put('test_df',df_final_test, format='table', data_colum
ns=True)
hdf.close()

#=====
=====

```

2. SVD svd_dot feature

```

In [44]: # return the dot product of U,V.T --> SVD feature
def dot(a,b):
    return np.dot(a,b)

```

2.1 SVD Dot product of U,V.T on source node and destination node

```

In [46]: df_final_train = read_hdf('data/fea_sample/storage_sample_stage
5.h5', 'train_df',mode='r')
df_final_test = read_hdf('data/fea_sample/storage_sample_stage
5.h5', 'test_df',mode='r')

if not os.path.isfile('data/fea_sample/storage_sample_stage6.h
5'):
    #=====
    =====

    # applying dot product of U,V.T on source node and destinat
ion node

    df_final_train['source_U.VT']=df_final_train.apply(lambda r
ow: dot(row[['svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',

'svd_u_s_4','svd_u_s_5', 'svd_u_s_6']],

row[['svd_v_s_1', 'svd_v_s_2','svd_v_s_3',

'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6']]),axis=1)

    df_final_test['source_U.VT']=df_final_test.apply(lambda row
: dot(row[['svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',

'svd_u_s_4','svd_u_s_5', 'svd_u_s_6']],

row[['svd_v_s_1', 'svd_v_s_2','svd_v_s_3',

```

```

'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6']]),axis=1)

df_final_train["destination_U.VT"] = df_final_train.apply(l
ambda row: dot(row[['svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',

'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6']],

row[['svd_v_d_1','svd_v_d_2', 'svd_v_d_3',

'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6']]),axis=1)

df_final_test["destination_U.VT"] = df_final_test.apply(lam
bda row: dot(row[['svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',

'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6']],

row[['svd_v_d_1','svd_v_d_2', 'svd_v_d_3',

'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6']]),axis=1)

hdf = HDFStore('data/fea_sample/storage_sample_stage6.h5')
hdf.put('train_df',df_final_train, format='table', data_col
umns=True)
hdf.put('test_df',df_final_test, format='table', data_colum
ns=True)
hdf.close()

#=====
=====

```

2.2 SVD Dot product of U,U.T and V,V.T on source and destination node

```

In [47]: df_final_train = read_hdf('data/fea_sample/storage_sample_stage
6.h5', 'train_df',mode='r')
df_final_test = read_hdf('data/fea_sample/storage_sample_stage
6.h5', 'test_df',mode='r')

if not os.path.isfile('data/fea_sample/storage_sample_stage7.h
5'):
    #=====
    =====

    # applying Dot product of U,U.T on source node and V,V.T on
    destination node

    df_final_train['source_destination_U.UT']=df_final_train.ap
    ply(lambda row: dot(row[['svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',

'svd_u_s_4','svd_u_s_5', 'svd_u_s_6']],

row[['svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',

'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6']]),axis=1)

    df_final_test['source_destination_U.UT']=df_final_test.appl
    y(lambda row: dot(row[['svd_u_s_1', 'svd_u_s_2', 'svd_u_s_3',

```

```

'svd_u_s_4', 'svd_u_s_5', 'svd_u_s_6']],

row[['svd_u_d_1', 'svd_u_d_2', 'svd_u_d_3',

'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6']]),axis=1)

df_final_train["source_destination_V.VT"] = df_final_train.
apply(lambda row: dot(row[['svd_v_s_1', 'svd_v_s_2','svd_v_s_3'
,

'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6']]),

row[['svd_v_d_1','svd_v_d_2', 'svd_v_d_3',

'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6']]),axis=1)

df_final_test["source_destination_V.VT"] = df_final_test.ap
ply(lambda row: dot(row[['svd_v_s_1', 'svd_v_s_2','svd_v_s_3',

'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_6']]),

row[['svd_v_d_1','svd_v_d_2', 'svd_v_d_3',

'svd_v_d_4', 'svd_v_d_5', 'svd_v_d_6']]),axis=1)
hdf = HDFStore('data/fea_sample/storage_sample_stage7.h5')
hdf.put('train_df',df_final_train, format='table', data_col
umns=True)
hdf.put('test_df',df_final_test, format='table', data_colum
ns=True)
hdf.close()

#=====
=====

```

3. Modeling

```

In [2]: #reading
from pandas import read_hdf
df_final_train = read_hdf('data/fea_sample/storage_sample_stage
7.h5', 'train_df',mode='r')
df_final_test = read_hdf('data/fea_sample/storage_sample_stage
7.h5', 'test_df',mode='r')

```

```

In [3]: df_final_train.columns

```

```

Index(['source_node', 'destination_node', 'indicator_li
nk',
      'jaccard_followers', 'jaccard_followees', 'cosin
e_followers',
      'cosine_followees', 'num_followers_s', 'num_foll
owees_s',
      'num_followees_d', 'inter_followers', 'inter_fol
lowees', 'adar_index',
      'follows_back', 'same_comp', 'shortest_path', 'w
eight_in', 'weight_out',
      'weight_f1', 'weight_f2', 'weight_f3', 'weight_f

```

```

4', 'page_rank_s',
    'page_rank_d', 'katz_s', 'katz_d', 'hubs_s', 'hu
bs_d', 'authorities_s',
    'authorities_d', 'svd_u_s_1', 'svd_u_s_2', 'svd_
u_s_3', 'svd_u_s_4',
    'svd_u_s_5', 'svd_u_s_6', 'svd_u_d_1', 'svd_u_d_
2', 'svd_u_d_3',
    'svd_u_d_4', 'svd_u_d_5', 'svd_u_d_6', 'svd_v_s_
1', 'svd_v_s_2',
    'svd_v_s_3', 'svd_v_s_4', 'svd_v_s_5', 'svd_v_s_
6', 'svd_v_d_1',
    'svd_v_d_2', 'svd_v_d_3', 'svd_v_d_4', 'svd_v_d_
5', 'svd_v_d_6',
    'preferential_score_source', 'preferential_score
_destination',
    'preferential_score_source_destination', 'source
_U.VT',
    'destination_U.VT', 'source_destination_U.UT',
    'source_destination_V.VT'],
dtype='object')

```

```
In [4]: df_final_train.shape
```

```
(100002, 61)
```

```
In [5]: df_final_train.head()
```

	source_node	destination_node	indicator_link	jaccard_followers
0	273084	1505602	1	0
1	832016	1543415	1	0
2	1325247	760242	1	0
3	1368400	1006992	1	0
4	140165	1708748	1	0

```
5 rows × 61 columns
```

```
In [6]: y_train = df_final_train.indicator_link
y_test = df_final_test.indicator_link
```

```
In [7]: df_final_train.drop(['source_node', 'destination_node', 'indicat
or_link'], axis=1, inplace=True)
df_final_test.drop(['source_node', 'destination_node', 'indicato
r_link'], axis=1, inplace=True)
```

3.1 RandomForestClassifier using various n_estimator parameter

```
In [54]: estimators = [10, 50, 100, 250, 450]
train_scores = []
test_scores = []
```

```

for i in estimators:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None,
                                criterion='gini',
                                max_depth=5, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=52, min_samples_split=120,
                                min_weight_fraction_leaf=0.0, n_estimators=i, n_jobs=-1,
                                random_state=25, verbose=0, warm_start=False)
    clf.fit(df_final_train, y_train)
    train_sc = f1_score(y_train, clf.predict(df_final_train))
    test_sc = f1_score(y_test, clf.predict(df_final_test))
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('Estimators = ', i, 'Train Score', train_sc, 'test Score', test_sc)
plt.plot(estimators, train_scores, label='Train Score')
plt.plot(estimators, test_scores, label='Test Score')
plt.xlabel('Estimators')
plt.ylabel('Score')
plt.title('Estimators vs score at depth of 5')


```

```

Estimators = 10 Train Score 0.9160210656883666 test Score 0.8986979722518678
Estimators = 50 Train Score 0.9236173001310616 test Score 0.9206015353875276
Estimators = 100 Train Score 0.9239773429238517 test Score 0.9204208934534687
Estimators = 250 Train Score 0.9227762915900352 test Score 0.9185402684563759
Estimators = 450 Train Score 0.9221044726301736 test Score 0.917974471646788

```

Text(0.5, 1.0, 'Estimators vs score at depth of 5')




```

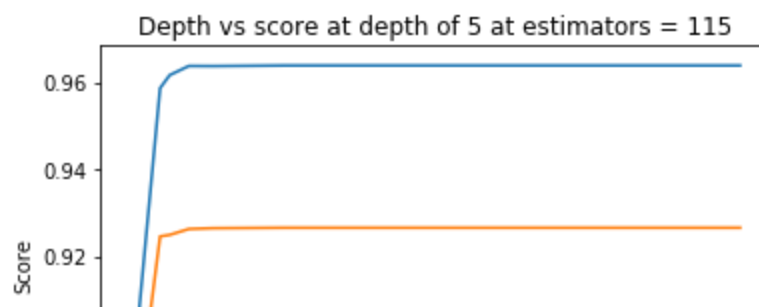
In [55]: depths = [3,9,11,15,20,35,50,70,130]
         train_scores = []
         test_scores = []
         for i in depths:
             clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                         max_depth=i, max_features='auto', max_leaf_nodes=None,
                                         min_impurity_decrease=0.0, min_impurity_split=None,
                                         min_samples_leaf=52, min_samples_split=120,
                                         min_weight_fraction_leaf=0.0, n_estimators=115, n_jobs=-1, random_state=25, verbose=0, warm_start=False)
             clf.fit(df_final_train,y_train)
             train_sc = f1_score(y_train,clf.predict(df_final_train))
             test_sc = f1_score(y_test,clf.predict(df_final_test))
             test_scores.append(test_sc)
             train_scores.append(train_sc)
             print('depth = ',i,'Train Score',train_sc,'test Score',test_sc)
         plt.plot(depths,train_scores,label='Train Score')
         plt.plot(depths,test_scores,label='Test Score')
         plt.xlabel('Depth')
         plt.ylabel('Score')
         plt.title('Depth vs score at depth of 5 at estimators = 115')
         plt.show()


```

```

depth = 3 Train Score 0.8885341622734241 test Score 0.8729199472273056
depth = 9 Train Score 0.9586056644880175 test Score 0.9245207583696314
depth = 11 Train Score 0.9615954230995517 test Score 0.9248649671087199
depth = 15 Train Score 0.9637023225649549 test Score 0.9262246870010924
depth = 20 Train Score 0.963679995139782 test Score 0.9264041994750656
depth = 35 Train Score 0.963817854178056 test Score 0.9265067069714719
depth = 50 Train Score 0.963817854178056 test Score 0.9265067069714719
depth = 70 Train Score 0.963817854178056 test Score 0.9265067069714719
depth = 130 Train Score 0.963817854178056 test Score 0.9265067069714719

```





3.2 RandomForestClassifier hypertuning using RandonSearchCV

```
In [56]: from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

param_dist = {"n_estimators":sp_randint(105,125),
              "max_depth": sp_randint(10,15),
              "min_samples_split": sp_randint(110,190),
              "min_samples_leaf": sp_randint(25,65)}

clf = RandomForestClassifier(random_state=25,n_jobs=-1)

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
                               n_iter=5,cv=10,scoring='f1',
                               random_state=25, return_train_score=True)

rf_random.fit(df_final_train,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
print('mean train scores',rf_random.cv_results_['mean_train_score'])
```

```
mean test scores [0.96241759 0.96203769 0.9607686  0.961
76077 0.96259747]
mean train scores [0.96287389 0.96265065 0.96118371 0.96
245331 0.9634601 ]
```

```
In [57]: rf.random.best_estimator
```

```
RandomForestClassifier(bootstrap=True, class_weight=None,
                       criterion='gini',
                       max_depth=14, max_features='auto',
                       max_leaf_nodes=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=28, min_samples_split=111,
                       min_weight_fraction_leaf=0.0, n_estimators=121,
                       n_jobs=-1, oob_score=False, random_state=None)
```

```
om_state=25, verbose=0,  
warm_start=False)
```

```
In [58]: #best hyperparameter after hypertuning  
rf_random.best_params_
```

```
{'max_depth': 14,  
 'min_samples_leaf': 28,  
 'min_samples_split': 111,  
 'n_estimators': 121}
```

```
In [59]: # best score  
rf_random.best_score_
```

```
0.9625974691596553
```

3.2.1 Training using best hyperparameter

```
In [8]: # training using best hyperparameter  
clf = RandomForestClassifier(bootstrap=True, class_weight=None,  
criterion='gini',  
                                max_depth=14, max_features='auto', max_l  
eaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_  
split=None,  
                                min_samples_leaf=28, min_samples_split=1  
11,  
                                min_weight_fraction_leaf=0.0, n_estimato  
rs=121,  
                                n_jobs=-1, oob_score=False, random_state  
=25, verbose=0,  
                                warm_start=False)
```

```
In [9]: clf.fit(df_final_train,y_train)  
y_train_pred = clf.predict(df_final_train)  
y_test_pred = clf.predict(df_final_test)
```

```
In [10]: from sklearn.metrics import f1_score  
print('Train f1 score',f1_score(y_train,y_train_pred))  
print('Test f1 score',f1_score(y_test,y_test_pred))
```

```
Train f1 score 0.9644600149855206
```

```
Test f1 score 0.9267012408927708
```

```
In [11]: # utility function  
from sklearn.metrics import confusion_matrix  
def plot_confusion_matrix(test_y, predict_y):  
    C = confusion_matrix(test_y, predict_y)  
  
    A = (((C.T) / (C.sum(axis=1))) .T)  
  
    B = (C / C.sum(axis=0))  
    plt.figure(figsize=(20,4))  
  
    labels = [0,1]  
    # representing A in heatmap format
```

```

cmap=sns.light_palette("blue")
plt.subplot(1, 3, 1)
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", annot_kws=
{"size": 15}, xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f",annot_kws={
"size": 15}, xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt='.2f', annot_kws=
{"size": 15}, xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()


```

```


In [12]: print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)

```

Train confusion_matrix



Test confusion_matrix




```


In [13]: from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f
)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')

```

```
plt.legend()
plt.show()
```



```
In [14]: features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r',
align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



3.3 Xgboost hypertuning using

RandonSearchCV

```
In [74]: from sklearn.metrics import f1_score
         from sklearn.metrics import f1_score
         import xgboost as xgb
         from sklearn.model_selection import RandomizedSearchCV
         from scipy.stats import randint as sp_randint
         from scipy.stats import uniform

         param_dist = {"n_estimators": [10, 25, 50, 100, 150],
                        "max_depth": [3, 5, 7, 10, 15, 20, 25],
                        "learning_rate": [0.01, 0.05, 0.1, 0.3, 0.5]}

         clf = xgb.XGBClassifier(random_state=25, n_jobs=-1)

         gscv_xgboost = RandomizedSearchCV(clf, param_distributions=param_dist,
                                           n_iter=5, cv=10, scoring='f1',
                                           random_state=20, return_train_score=True)

         gscv_xgboost.fit(df_final_train, y_train)

         print('mean test scores', gscv_xgboost.cv_results_['mean_test_score'])
         print('mean train scores', gscv_xgboost.cv_results_['mean_train_score'])
```

```
mean test scores [0.97275645 0.97191409 0.96317531 0.973
69547 0.97616364]
mean train scores [0.97361705 0.97236247 0.96455085 0.97
454971 0.99949719]
```

```
In [75]: gscv_xgboost.best_estimator_
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0,
              learning_rate=0.3, max_delta_step=0, max_depth=25,
              min_child_weight=1, missing=None, n_estimators=10, n_jobs=-1,
              nthread=None, objective='binary:logistic', random_state=25,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=None, subsample=1, verbosity=1)
```

```
In [78]: #best hyperparameter after hypertuning
         gscv_xgboost.best_params_
```

```
{'n_estimators': 10, 'max_depth': 25, 'learning_rate': 0.3}
```

```
In [79]: # best score
```

```
gscv_xgboost.best_score_
```

0.9761636355560456

3.3.1 Training using best hyperparameter

```
In [15]: # training using best hyperparameter
clf = xgb.XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0,
                        learning_rate=0.3, max_delta_step=0, max_depth=25,
                        min_child_weight=1, missing=None, n_estimators=10,
                        n_jobs=-1, nthread=None, objective='binary:logistic', random_state=25,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                        silent=None, subsample=1, verbosity=1)
```

```
In [16]: clf.fit(df_final_train,y_train)
y_train_pred = clf.predict(df_final_train)
y_test_pred = clf.predict(df_final_test)
```


```
In [17]: from sklearn.metrics import f1_score
print('Train f1 score',f1_score(y_train,y_train_pred))
print('Test f1 score',f1_score(y_test,y_test_pred))
```

Train f1 score 0.9995604922486815


Test f1 score 0.9331087482892935

```
In [18]: print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)
```


Train confusion_matrix




Test confusion_matrix



```
In [21]: from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()
plt.show()
```



```
In [22]: features = df_final_train.columns
importances = clf.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r',
align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```

4. Conclusion

```
In [58]: from prettytable import PrettyTable
# After adding preferential_attachment and SVDdot feature
print("====Before adding preferential_attachment and SVDdot
feature====")
x=PrettyTable()
x.field_names = ["Model","train_f1_score","test_f1_score"]
x.add_row(["Random forest",0.965,0.924])
print(x)
print("="*70, "\n")

# After adding preferential_attachment and SVDdot feature
print("====After adding preferential_attachment and SVDdot f
eature====")
x=PrettyTable()
x.field_names = ["Model","train_f1_score","test_f1_score"]
x.add_row(["Random forest",0.964,0.926])
x.add_row(["Xgboost",0.999, 0.933])
print(x)
print("="*70)
```

```
====Before adding preferential_attachment and SVDdot
feature=====
```

```
+-----+-----+-----+
|      Model      | train_f1_score | test_f1_score |
+-----+-----+-----+
```

```
| Random forest |      0.965      |      0.924      |
+-----+-----+-----+
```

```
=====
```

```
=====After adding preferential_attachment and SVDdot feature=====
```

Model	train_f1_score	test_f1_score
Random forest	0.963	0.926
Xgboost	0.99	0.932

```
=====
```

1. Although difference in result in RandomForest model is not much after adding preferential_attachment and SVDdot feature, but feature importance of model is clearly showing preferential_attachment is an important feature
1. Using Xgboost with hyperparameter tuning is giving best result among all.

```
| test_f1_score = 0.933
```

END:)